


Network Architecture


- Network Architecture: framework for designing and implementing Networks
- Components:
 - Software
 - Protocols
 - Services
 - Hardware
 - Transmission technology, media and devices
 - Scale: LANs, MANs, and WANs
 - Topology

Subnet



Layered Architecture

- The purpose of each layer is to offer a communication services to higher level layers
- Each layer has two interfaces
 - peer-to-peer interface
 - defines the form and types of messages exchanged between peers (indirect communication)
 - service interface
 - defines the primitives (operations) that a layer provides to the layer above it
- Layering is non-linear




Protocols

- The functionality encapsulated within each layer is called *Protocol*
- The Protocol refers to both
 - the abstract peer-to-peer & service interfaces and
 - the objects that implement those interfaces
- Protocol vs. Service
 - Service is the set of primitives provided to the higher layer
 - Protocol defines the implementation of these primitives
- Protocol stack

ISO OSI Architecture

- International Standards Organisation (ISO)
- *Physical*: transmission of raw bits onto the communications medium
- *Data link*: reliable transmission of frames, flow control, arbitration
- *Network*: packet switching, routing congestion control
- *Transport*: process-to-process channel, node-to-node connection, provides user services, flow control, multiplexing
- *Session*
- *Presentation*
- *Application*


Internet Architecture



- Host-to-Network Layer (OSI Physical and Data link layers)
- Internet Layer (OSI Network layer - Internet Protocol/IP)
- **Transport Layer**
 - **Transmission Control Protocol - TCP**
 - **User Datagram Protocol - UDP**
- Application Layer

Accessing Transport Layer

- How application programs use protocol software (invoke the services) of transport layer to communicate across networks




Inter-process Communication

- Message passing model: send - receive
- Synchronous communications : blocking
- Asynchronous communications
 - Blocking (receive) vs non blocking (send, receive)
- Reliability
 - Validity: messages are guaranteed to be delivered despite of packets being lost
 - Integrity: messages must arrive uncorrupted and without duplication
- Ordering: message send and delivery

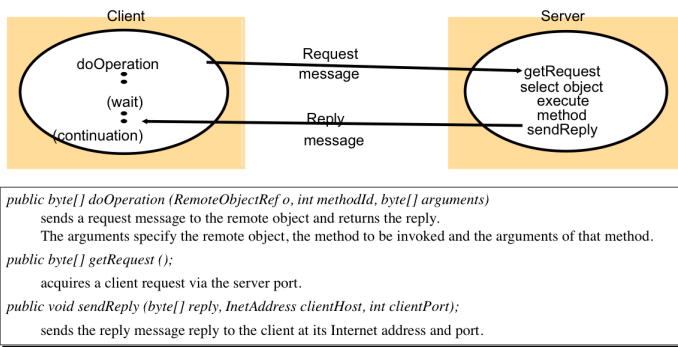
Sockets and ports

- Sockets provide an abstract endpoint for inter-process communication (UDP, TCP)
- Destination:
 - internet address, port, process id
- Port: message destination within a computer (integer) –Internet addresses: names rather than numbers (2^{16})
- A socket must be bound to a local port and an internet address
- A process may use multiple ports but cannot share ports (exception: multicast)



The Request-Reply Protocol

- Processes communicate using a protocol



Issues

- Message identifiers (sequence+sender id)
- Failure Model
- Timeouts
- Duplicate request messages
- Lost reply messages
- History for retransmissions

UDP Sockets

- Receive method returns the IP address+port of sender so that a reply can be sent
- Message size: receiving process specifies a byte-array in which to receive the message (IP: 2^{16} bytes, usually 8kB)
- Non-blocking send & blocking receive (timeouts and threads for deadlocks)
- Receive *from any*

Java API

- `DatagramPacket` class provides
 - a constructor that makes an instance of a message: [message byte array, length, IP address, port no]
 - A constructor for use when receiving a message
 - Methods: `getData`, `getPort`, `getAddress`
- `DatagramSocket` class provides
 - A constructor that takes a port as an argument for use by processes that need to use a port
 - A no-argument constructor that allows the system to choose a free local port
 - `SocketException` if problem with port
 - Methods:
 - `send`: argument an instance of `DatagramPacket`
 - `Receive`: argument an empty `DatagramPacket`
 - `setSoTimeout`: `InterruptedException` when expire
 - `connect`: connect a socket to a particular remote port and IP address, the socket is able to send to and receive from that address

TCP Sockets

- API provides an abstraction of a stream of bytes to which data may be written and from which data may be read. The abstraction hides the following:
 - Message sizes. The application can choose how much data it writes to a stream. TCP may fragment it to smaller packets
 - Lost messages
 - Flow control
 - Message duplication and ordering
 - Message destinations: once a connection is established, processes simply read from and write to a stream (no need for IP addresses and ports)
- Client-server model during connection:
 - client creates a stream socket bound to a port and asks for a connection to a server port
 - server creates a listening socket bound a a port and waits to *accept* connect requests
- During operation each socket is both for input and output
- Close a socket when no more data to write

Java API

- *ServerSocket* class: create a socket at a server port for listening for connect requests
 - Method *accept*: wait until there is a *connect* request in the queue, then create an instance of *Socket*
- *Socket* class: client uses a constructor which creates a socket (specifying the DNS hostname and port of a server) and connects it to the specified remote server (*UnknownHostException*, *IOException*)
 - Methods: *getInputStream* and *getOutputStream*
 - Return types are abstract classes that define methods for reading (*InputStream*) and writing (*OutputStream*) bytes
 - Return values can be used as the arguments of constructors for suitable input and output streams.

Group Communication

- Multicast: an operation that sends a single message from one process to each of the members of a group of processes
- Fault tolerance based on replicated services
- Finding the discovery servers in spontaneous networking
- Better performance through replicated data
- Propagation of event notifications

IP Multicast

- multicast group is specified by a class D Internet address
 - membership is dynamic, to join make a socket
 - programs using multicast use UDP and send datagrams to multicast addresses and (ordinary) port
 - (For example of Java code see book)